# Model Selection and Hypothesis Testing

732A90
Computational Statistics

Krzysztof Bartoszek
(krzysztof.bartoszek@liu.se)

II 2019 ()
Department of Computer and Information Science
Linköping University

Tools for model selection

- **Comparing different models**
- Information criteria (not this course)
- Cross–validation
- Hypothesis testing

- **Uncertainty estimation**
- Confidence intervals

Hypothesis testing: **Recap**

1. Assume a probabilistic model
   State a null hypothesis (H_0 e.g. no difference) and
   alternative (H_1 difference)

2. Observe data X

3. Calculate a test statistic e.g. $T(X) = (\overline{X})/(\widehat{sd(X)})$
   (different statistics will have different **efficiency** (power,
   ability to distinguish between hypotheses) associated with
   them)

4. Under H_0 $T(X)$ has "known" distribution

5. Decision: Is the value of $T(X)$ *surprising* (in the **critical
   region**)? If so reject H_0 in favour of H_1.

x<–**rnorm**(10,**mean**=4,**sd**=1)

Hypotheses:

$H_0 : \mu = 4, X \sim \mathcal{N}(\mu, \sigma^2)$
$H_1 : \mu \neq 4, X \sim \mathcal{N}(\mu, \sigma^2)$

Hypothesis testing: Example

```
x<-rnorm(10,mean=4,sd=1)
```

Hypotheses:

$$H_0 : \mu = 4, X \sim \mathcal{N}(\mu, \sigma^2)$$
$$H_1 : \mu \neq 4, X \sim \mathcal{N}(\mu, \sigma^2)$$

Test statistic

$$T(x) = \frac{\bar{x} - \mu}{s/\sqrt{n}} \sim t(n-1)$$

```
tx<-(mean(x)-4)/(sqrt(var(x)/length(x)))
t0<-qt(0.975,df=length(x)-1)
(tx>t0)||(tx<(-t0)) ## reject if TRUE
```

## Hypothesis testing: Example

x<−**rnorm**(10,**mean**=4,**sd**=1)

Hypotheses:

$H_0 : \mu = 4, X \sim \mathcal{N}(\mu, \sigma^2)$
$H_1 : \mu \neq 4, X \sim \mathcal{N}(\mu, \sigma^2)$

Test statistic

$T(x) = \frac{\bar{x} - \mu}{s/\sqrt{n}} \sim t(n-1)$

tx<−(**mean**(x)−4)/(**sqrt**(**var**(x)/**length**(x)))
t0<−**qt**(0.975,**df**=**length**(x)−1)
(tx>t0)||(tx<(−t0)) *## reject if TRUE*

How does one compares different statistics?
**POWER**
Power $= 1 -$ Type II error

Ability to correctly identify *surprise*,
i.e. indicate H_1.

How to compute power?

- Analytically (?)

- Generate data samples that satisfy H_1
  Compute percent of correct rejections

We may use "any" test statistic.
We do **not** need to know its distribution.

$H_0 : \mu = 4, X \sim \mathcal{N}(\mu, \sigma^2)$
$H_1 : \mu \neq 4, X \sim \mathcal{N}(\mu, \sigma^2)$

Monte Carlo Hypothesis testing

We may use "any" test statistic.
We do **not** need to know its distribution.

$H_0 : \mu = 4, X \sim \mathcal{N}(\mu, \sigma^2)$
$H_1 : \mu \neq 4, X \sim \mathcal{N}(\mu, \sigma^2)$

Test statistic

$$T(x) = \frac{\bar{x} - \mu}{s/\sqrt{n}} \sim t(n-1)$$

1: **for** $i = 1$ to B **do**
2:   Generate $Y_1, \ldots, Y_n$ i.i.d. from H_0, i.e. $\mathcal{N}(\mathbf{4}, \sigma^2)$
3:   Compute t_i from $Y_1, \ldots, Y_n$
4: **end for**
5: Use $t_1, \ldots, t_B$ to construct a histogram
6: Use the histogram as the distribution of $T(x)$ under H_0

Monte Carlo Hypothesis testing


```
x<-rnorm(10,4,1)
s<-var(x)
B<-10000
n<-length(x)
tsamp<-rep(NA,B)
for (i in 1:B){
    Y<-rnorm(n,4,s)
    tsamp[i]<-(mean(Y)-4)/(sd(Y)/sqrt(length(Y)))
}
hist(tsamp,breaks=50,col=gray(0.8),main="",xlab="t"
    ,ylab="",freq=FALSE,cex.axis=1.5,cex.lab=1.5)
```

Permutation tests

- A. k. a. randomization tests

- One solution if we do not know the distribution under H_0

- Computationally expensive

- Any sample size

- Two sample problem:
  - Population 1 distributed as F
  - Population 2 distributed as G
 - $H_0 : F = G$
    $H_1 : F \neq G$

```
> t(mouse)
      [,1] [,2] [,3] [,4] [,5] [,6] [,7] [,8] [,9] [,10] [,11] [,12]
Group "y"  "z"  "z"  "y"  "y"  "z"  "y"  "y"  "y"  "y"   "z"   "z"
Value " 10" " 16" " 23" " 27" " 31" " 38" " 40" " 46" " 50" " 52" " 94" " 99"
      [,13] [,14] [,15] [,16]
Group "y"   "z"   "y"   "z"
Value "104" "141" "146" "197"
```

Do the values differ significantly between control and treatment groups?

IDEA: If $F = G$ then **group label does not matter**

We may permute labels and still have a sample from F (or G)

Before permutation

After permutation

Test statistic:

$$T(X) = \text{mean}(\text{values}|\text{group} = z) - \text{mean}(\text{values}|\text{group} = y)$$

1: $T(X)$ value of statistic from observed data

2: Create permutations $g_1^*, \ldots, g_B^*$ of group variable
{If the number of permutations is too large, sample B
randomly **without** replacement. E.g. generate random
permutations and keep only unique ones.}

3: Evaluate test statistic on each permutation

4: Estimate p–value: $\hat{p} = \#\{T(X_{g_b^*}) \geq T(X)\}/B$
5: If test is two–sided: $\hat{p} = \#\{|T(X_{g_b^*})| \geq |T(X)|\}/B$

Do we reject the null?

```
B=1000
stat=numeric(B)
n=dim(mouse)[1]
for(b in 1:B){
  Gb=sample(mouse$Group, n)
  stat[b]=mean(mouse$Value[Gb=='z'])-mean(mouse$
      Value[Gb=='y'])
}
stat0=mean(mouse$Value[mouse$Group=='z'])-mean(
    mouse$Value[mouse$Group=='y'])
print(c(stat0,mean(stat>stat0)))
## [1] 30.63492   0.12700
```
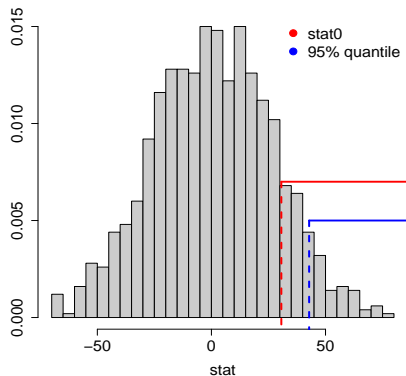
Resampling methods

Jackknife and bootstrap

Theory **different**, coding **similar**

Data (**i.i.d.**) $X \sim F(\cdot, w)$

1: Observed data: $D = (X_1, \ldots, X_n)$, estimator $\widehat{w} = T(D)$
2: **for** $i = 1, \ldots, B$ { Jackknife $B \leq n$} **do**
3:     Generate

       $D_i^* = (X_1^*, \ldots, X_n^*)$ by sampling with replacement
       {**Nonparametric Bootstrap**, F unknown}

       $D_i^* = X[-i]$ {**Jackknife**, F unknown}

       $D_i^* = (X_1^*, \ldots, X_n^*)$ by generating from $F(\cdot, \hat{w})$
       {**Parametric Bootstrap**, F known}

4: **end for**
5: Distribution of $\widehat{w}$ is estimated by $T(D_1^*), \ldots, T(D_B^*)$
    {The histogram based on resampled values is used in place
    of the true density.}

Uncertainty estimation: confidence intervals

Estimate $100(1 - \alpha)\%$ percentile confidence interval for w

$se(\cdot)$ is the square root of estimated variance (computationally heavy)

NOT by jackknife **TOO DEPENDENT!!**

1: Compute $T(D_1^*), \ldots, T(D_B^*)$
2: Sort in ascending order, obtaining $y_1, \ldots, y_B$
   {**percentile** method} **OR**
   Compute $y_i = (T(D_i^*) - T(D))/(se(T(D_i^*)))$ $i = 1, \ldots, B$
   {**t** method}
3: Define $A_1 = \lceil (B\alpha/2) \rceil$, $A_2 = \lfloor (B - B\alpha/2) \rfloor$
4: Confidence interval is given by
   (y_{A_1}, y_{A_2}) {**percentile** method} **OR**
   $(T(D) - se(T(D^*)) \cdot y_{A_1}, T(D) + se(T(D^*)) \cdot y_{A_2})$
   {**t** method}

Hypothesis testing: does statistic from observed data fall into
CI (H_0) or not (H_1)

Bootstrap

$$\widehat{\text{Var}\,[T(\cdot)]} = \frac{1}{B-1} \sum_{i=1}^{B} \left(T(D_i^*) - \overline{T(D^*)} \right)^2$$

Jackknife $(n = B)$

$$\widehat{\text{Var}\,[T(\cdot)]} = \frac{1}{n(n-1)} \sum_{i=1}^{n} \left(T_i^* \right) - J(T))^2,$$

where

$$T_i^* = nT(D) - (n-1)T(D_i^*) \quad J(T) = \frac{1}{n} \sum_{i=1}^{n} T_i^*$$

Bootstrap in R


```R
library("boot")
stat1<-function(data,vn){
    data<-as.data.frame(data[vn,])
    res<-lm(Response~Predictor,data)
    res$coefficients[2]
}
x<-rnorm(100);data<-cbind(Predictor=x,Response=3+2*
    x+rnorm(length(x),sd=0.5))
res<-boot(data,stat1,R=1000)
print(boot.ci(res))
## BOOTSTRAP CONFIDENCE INTERVAL CALCULATIONS
##Based on 1000 bootstrap replicates
#Intervals :
#Level         Normal               Basic
#95%    ( 1.933,  2.164 )     ( 1.935,  2.162 )
# Level        Percentile               BCa
#95%    ( 1.934,  2.161 )     ( 1.936,  2.166 )
```

Bootstrap bias correction

1: Observed data: $D = (X_1, \ldots, X_n)$, estimator $\widehat{w} = T(D)$
2: **for** $i = 1, \ldots, B$ **do**
3:     Generate

    $D_i^* = (X_1^*, \ldots, X_n^*)$ by sampling with replacement.
4:     Calculate $T_i^* = T(D_i^*)$.
5: **end for**
6: Bias corrected estimator is

$$T_1 := 2T(D) - \frac{1}{B} \sum_{i=1}^{B} T_i^*.$$

Jackknife also has a bias correction method (see 2016 slides).

Comments

- Jackknife overestimate variance

- Bootstrap–t method is more accurate than percentile

- Permutations: sampling **without** replacement, bootstrap **with**

- Permutation p–value exact if all permutations used, bootstrap always approximate

- Bootstrap may be used for a wider class of problems

- Nonparametric bootstrap works badly for small samples ($n < 40$)

- Parametric bootstrap can work for small samples

- Bias corrections

- Methods do not require distributional assumptions

Data predictors: `X[,c(V1,V2)]`, response: `Y`
Model M relating Y and X

Competing models
H_0 variables `V1` should not be in M (smaller model)
H_1 all variables are significant
Test statistic: $T(M)$

Permutation test

1: **for** $i = 1 \ldots B$ **do**
2: Obtain `V1`* by permuting order of columns in `V1`,
    fit model `Y=M(X[,c(V1*,V2)])`
3: Compute test statistic T_i for this model
4: **end for**
5: Compute p–value using above distribution of T

- Why are some models better than others?

- Hypothesis testing

- Monte Carlo hypothesis testing

- Resampling methods (permutations, jackknife, bootstrap)

- Simulation methods (parametric bootstrap)